



Balanced Trees


Sridhar Alagar

Motivation

Height of the BST can be as high $n-1$




Binary Search Tree



- In a BST, each node stores one key
- Each node splits the key space into two halves

Generalizing BST

- Each node stores many keys in sorted order
- Each node with k keys splits the key space into $k+1$ regions



Balanced Multiway Search Tree

- Building a balanced multiway search tree is easy
- Just pack them in one node

20	30	35	40	50	60	65	70
----	----	----	----	----	----	----	----

- At some point it becomes a bad idea. Why?
- It is just a sorted array
- Need to limit the number of keys in a node. How?

Strategy 1

- Push keys down after the limit exceeded in a node

20 30 70 80 35 40 50 38 39

Strategy 1

- Push keys down after the limit exceeded in a node

20

30 70 80 35 40 50 38 39

Strategy 1

- Push keys down after the limit exceeded in a node



70 80 35 40 50 38 39

Strategy 1

- Push keys down after the limit exceeded in a node

20	30	70
----	----	----

80 35 40 50 38 39

Strategy 1

- Push keys down after the limit exceeded in a node


20	30	70	80
----	----	----	----

Limit exceeded...
push down

35 40 50 38 39

Strategy 1


- Push keys down after the limit exceeded in a node



35 40 50 38 39

Strategy 1


- Push keys down after the limit exceeded in a node



40 50 38 39

Strategy 1


- Push keys down after the limit exceeded in a node



50 38 39

Strategy 1


- Push keys down after the limit exceeded in a node



38 39

Strategy 1


- Push keys down after the limit exceeded in a node



38 39


Strategy 1

- Push keys down after the limit exceeded in a node



Strategy 1

- Push keys down after the limit exceeded in a node



- Easy to implement
- But tree can be unbalanced

Strategy 2

- Split node and push middle key up if limit exceeded

20 30 70 80 35 40 50 38 39

Strategy 2

- Split node and push middle key up if limit exceeded

20	30	70
----	----	----

80 35 40 50 38 39

Strategy 2

- Split node and push middle key up if limit exceeded


20	30	70	80
----	----	----	----

Limit exceeded...
split and push up

35 40 50 38 39


Strategy 2

- Split node and push middle key up if limit exceeded




Strategy 2

- Split node and push middle key up if limit exceeded




Strategy 2

- Split node and push middle key up if limit exceeded




Strategy 2

- Split node and push middle key up if limit exceeded



Strategy 2


- Split node and push middle key up if limit exceeded



38 39 60 65

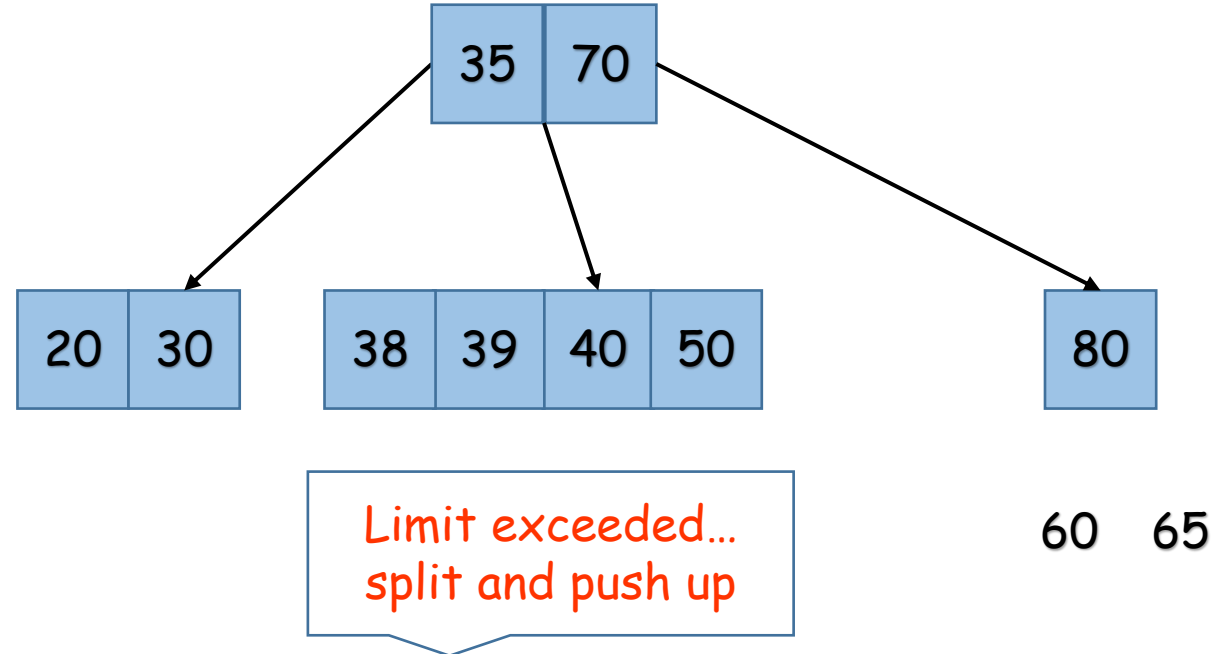
Strategy 2

- Split node and push middle key up if limit exceeded




Strategy 2

- Split node and push middle key up if limit exceeded




Strategy 2

- Split node and push middle key up if limit exceeded



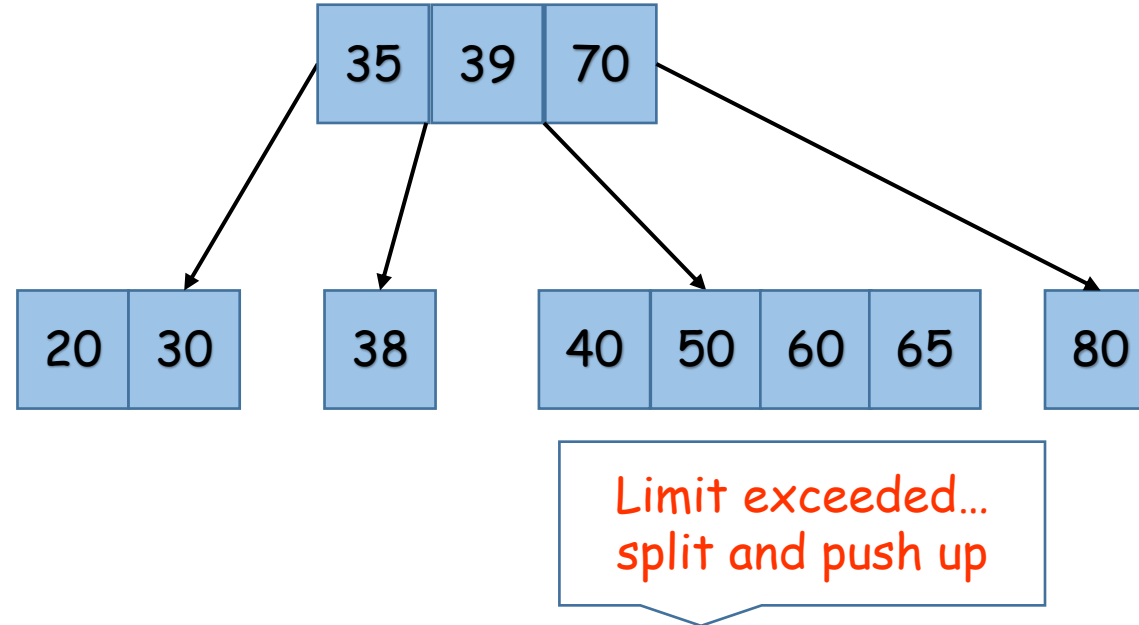
Strategy 2

- Split node and push middle key up if limit exceeded



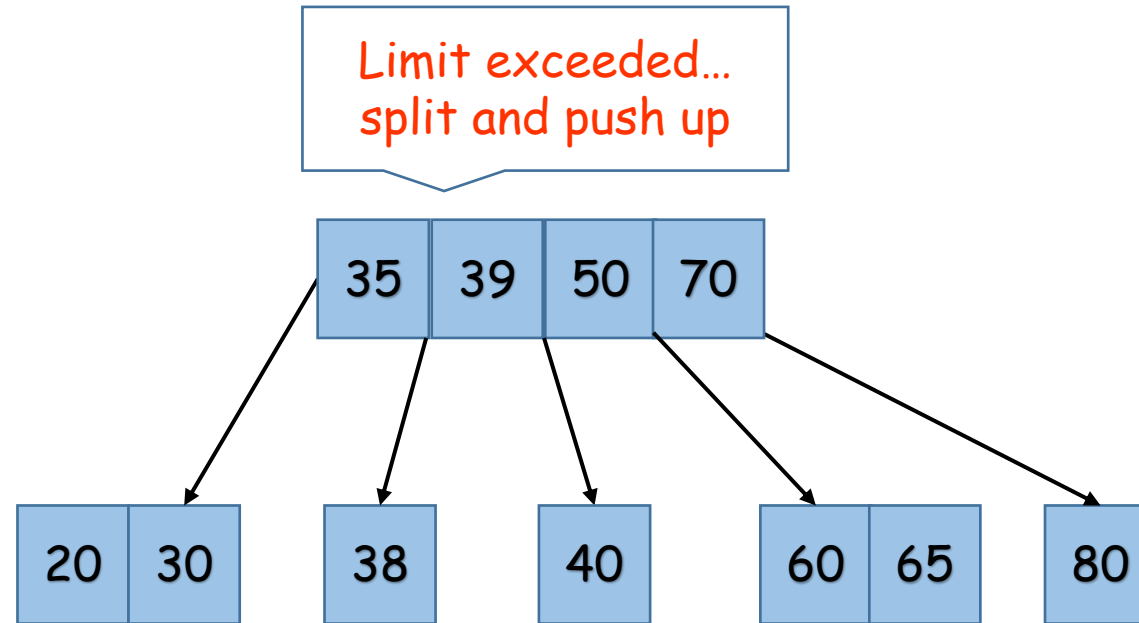
Strategy 2

- Split node and push middle key up if limit exceeded




Strategy 2

- Split node and push middle key up if limit exceeded



Strategy 2

- Split node and push middle key up if limit exceeded




- Slightly trickier to implement
- But tree is balanced. So, search, insert are efficient

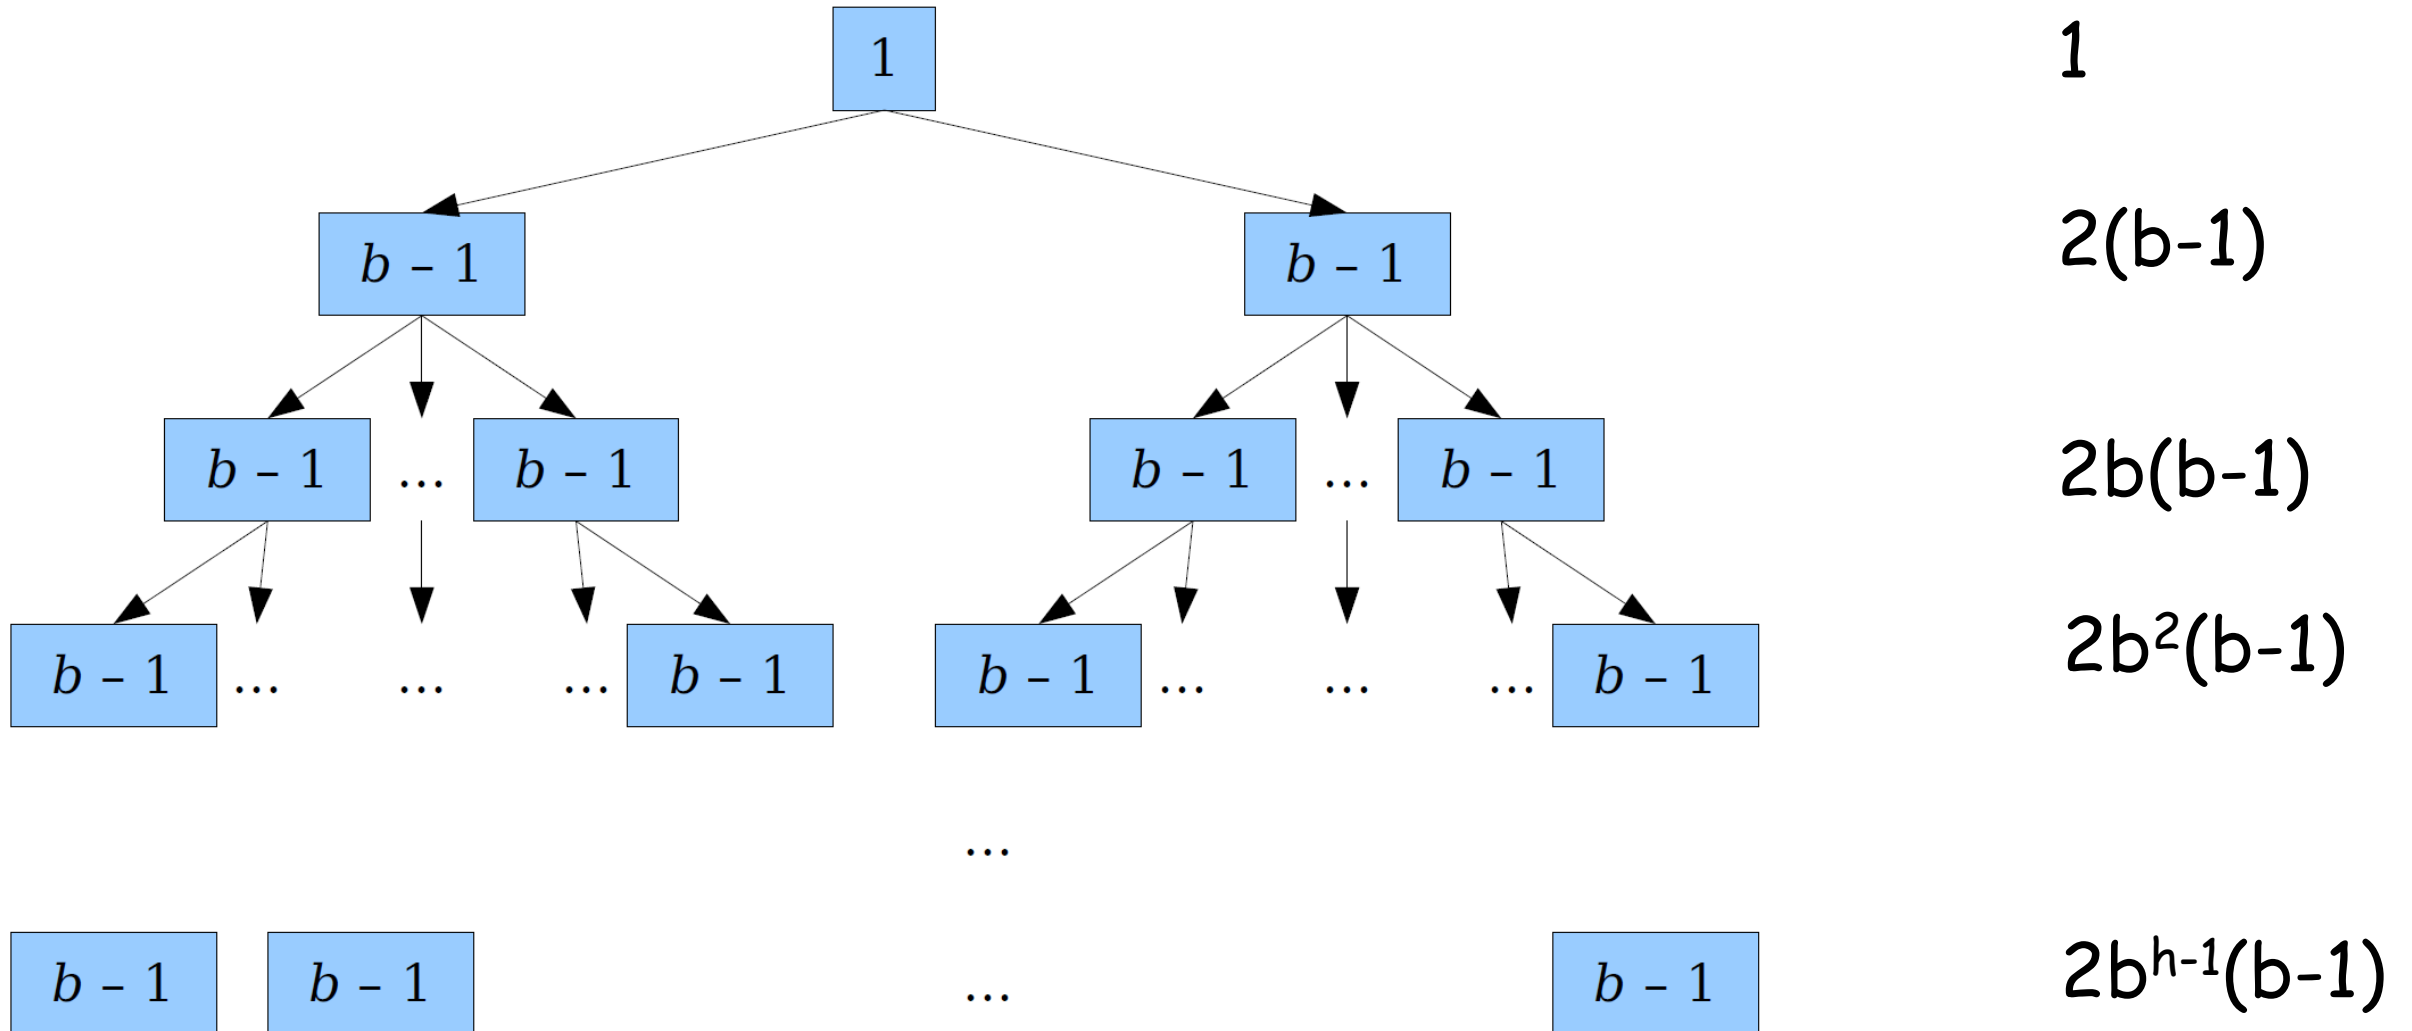
B-trees

A B-tree of order b is a multiway search tree where:

- each node has $b-1$ to $2b-1$ keys, other than root which can have up to $b-1$ keys
- each node is either a leaf or has one more child than the number of keys
- all leaves are at the same depth




Maximum height of a B-tree? $O(\log_b n)$




RT for Search

- Search every node
- Find the key or the child to descend to
- RT to search a node:
Binary search - $\log b$
- RT for search = $\log b \log_b n = O(\log n)$



RT for insertion

- Need to visit $\log_b n$ nodes
- May need to split every node $O(b)$
- RT for insertion: $= O(b \log_b n)$
 $= O((b/\log b) \log n)$
- In practice,
 - $b = 2$ (2-3-4 trees)
 - $b = 1024/4096$ (B-trees)



External data structures

- If n is too large to fit into memory, nodes are stored in secondary storage
- compute time \ll access time
- Ideal size (branching factor) of node is the size of block (minimum amount of data read from the storage device)
 - Typically, block size is 1K or 4K bytes
- Minimizes the storage device access

Insertion: $O(\log_b n)$


Search: $O(\log_b n)$

If there are billion keys, # disk access = 3

2-3-4 trees

A B-tree with order $b = 2$

- each node has 1 to 3 keys; root will have 1 key only
- each node is either a leaf or has one more child than the number of keys
- all leaves are at the same depth




The story so far

- We know how to build a balanced multiway search tree
- 2-3-4 trees when main memory is sufficient
- B-trees for huge amount of keys


Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree



Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree





39

Root Property: The root is black



Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree





Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree





Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree





Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree





Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree



Red Black Trees



A 2-3-4 tree can be transformed into a corresponding red-black tree



Root Property: The root is black
 Red Property: The parent of a red node is black



Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree





Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree





Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree



Red Black Trees

A 2-3-4 tree can be transformed into a corresponding red-black tree



What can you infer about the depth?

Red Black Trees

Root Property:


The root is black

Red Property:

The parent of a red node is black

Depth Property:

The number of black nodes on each path from the root to a leaf node is the same



Red Black Trees

Root Property:

The root is black

Red Property:


The parent of a red node is black

Depth Property:

The number of black nodes on each path from the root to a leaf node is the same



Leaf Nodes:

All leaf nodes are black nil nodes



Red Black Trees



A red-black tree can be transformed into a corresponding 2-3-4 tree



Red Black Trees

A red-black tree can be transformed into a corresponding 2-3-4 tree

Move the red nodes up to their parents




Corresponding Transformation


2-node




3-node




or




4-node



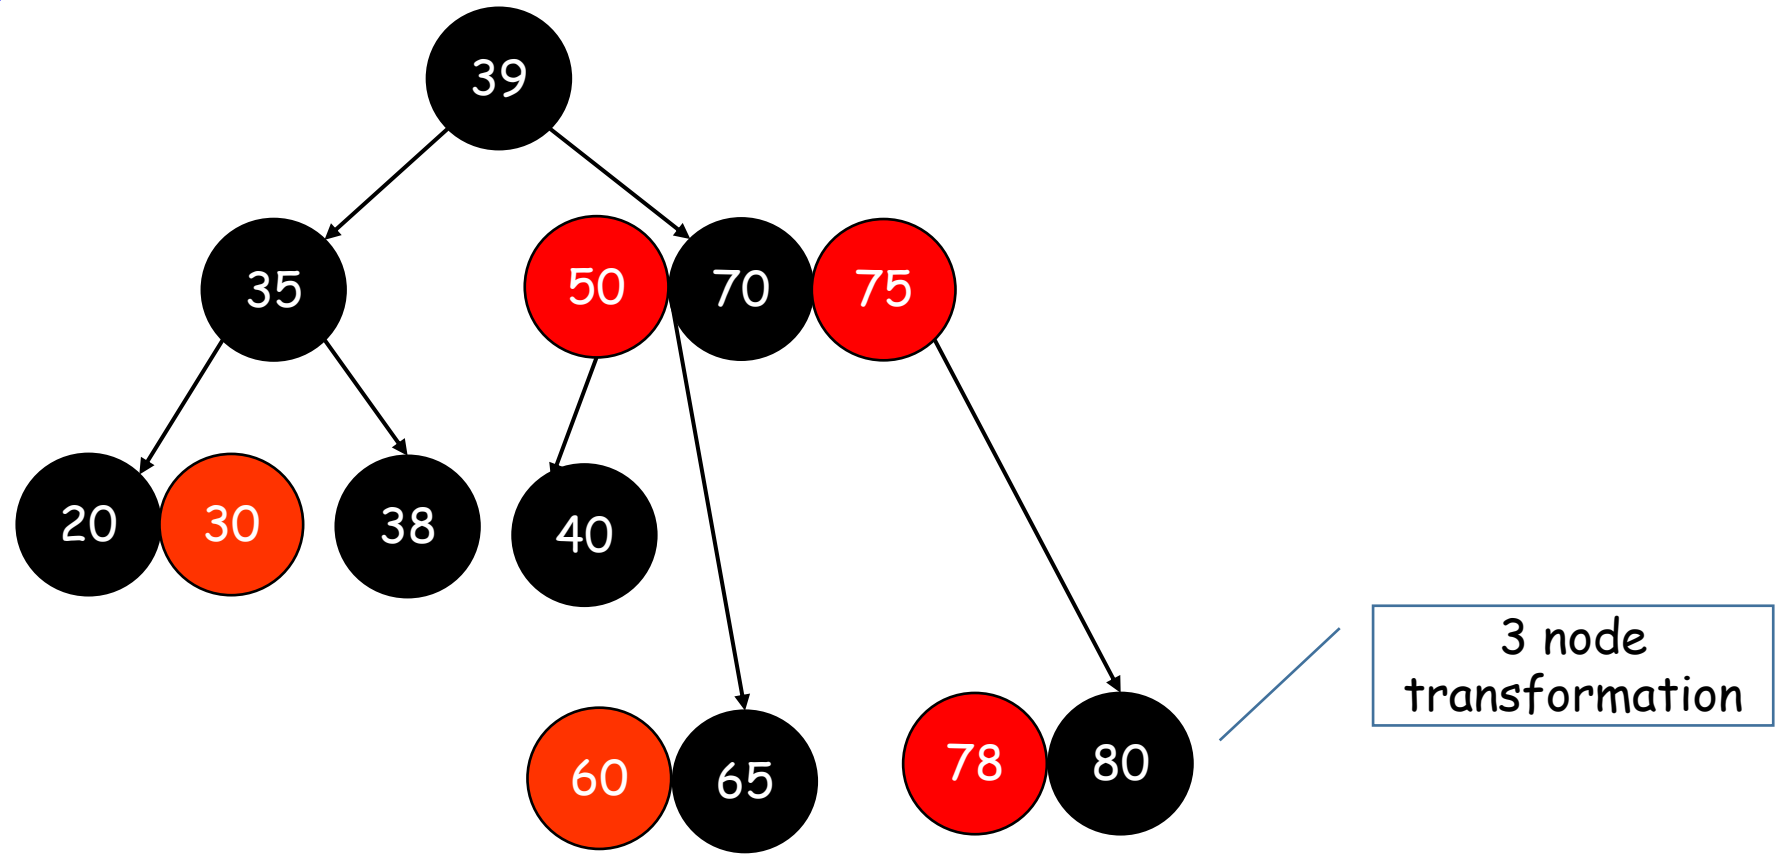
Insert(78)




Insert(78)




Insert(78)




Insert(78)




Insert(55)




Insert(55)




50	70	75
----	----	----




Insert(55)




50	70	75
----	----	----



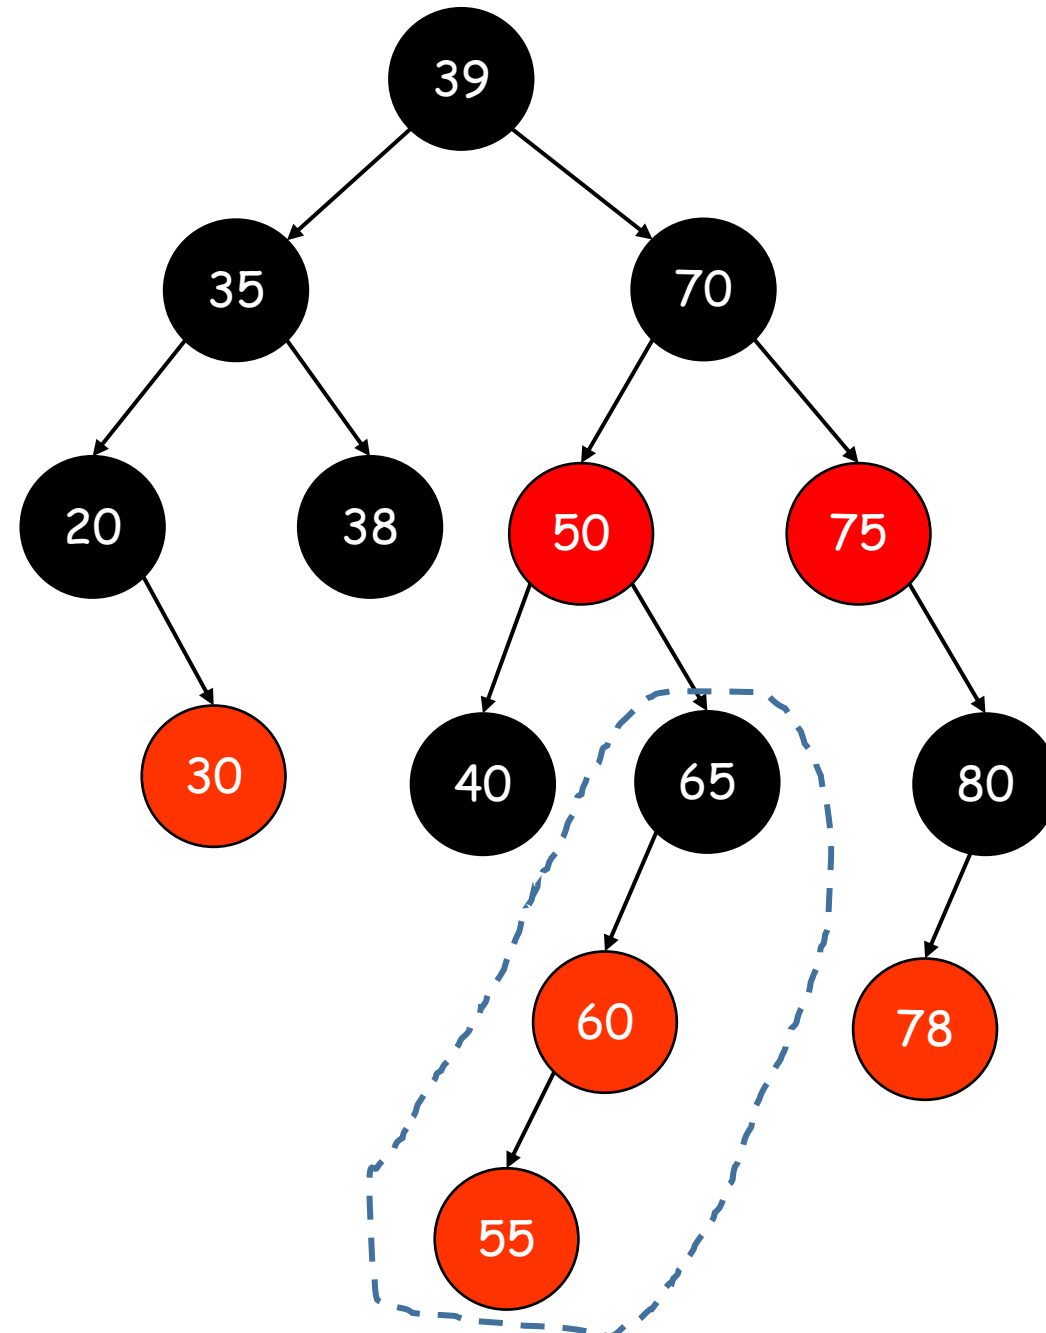
Insert(55)




Insert(55)





Insert(55)



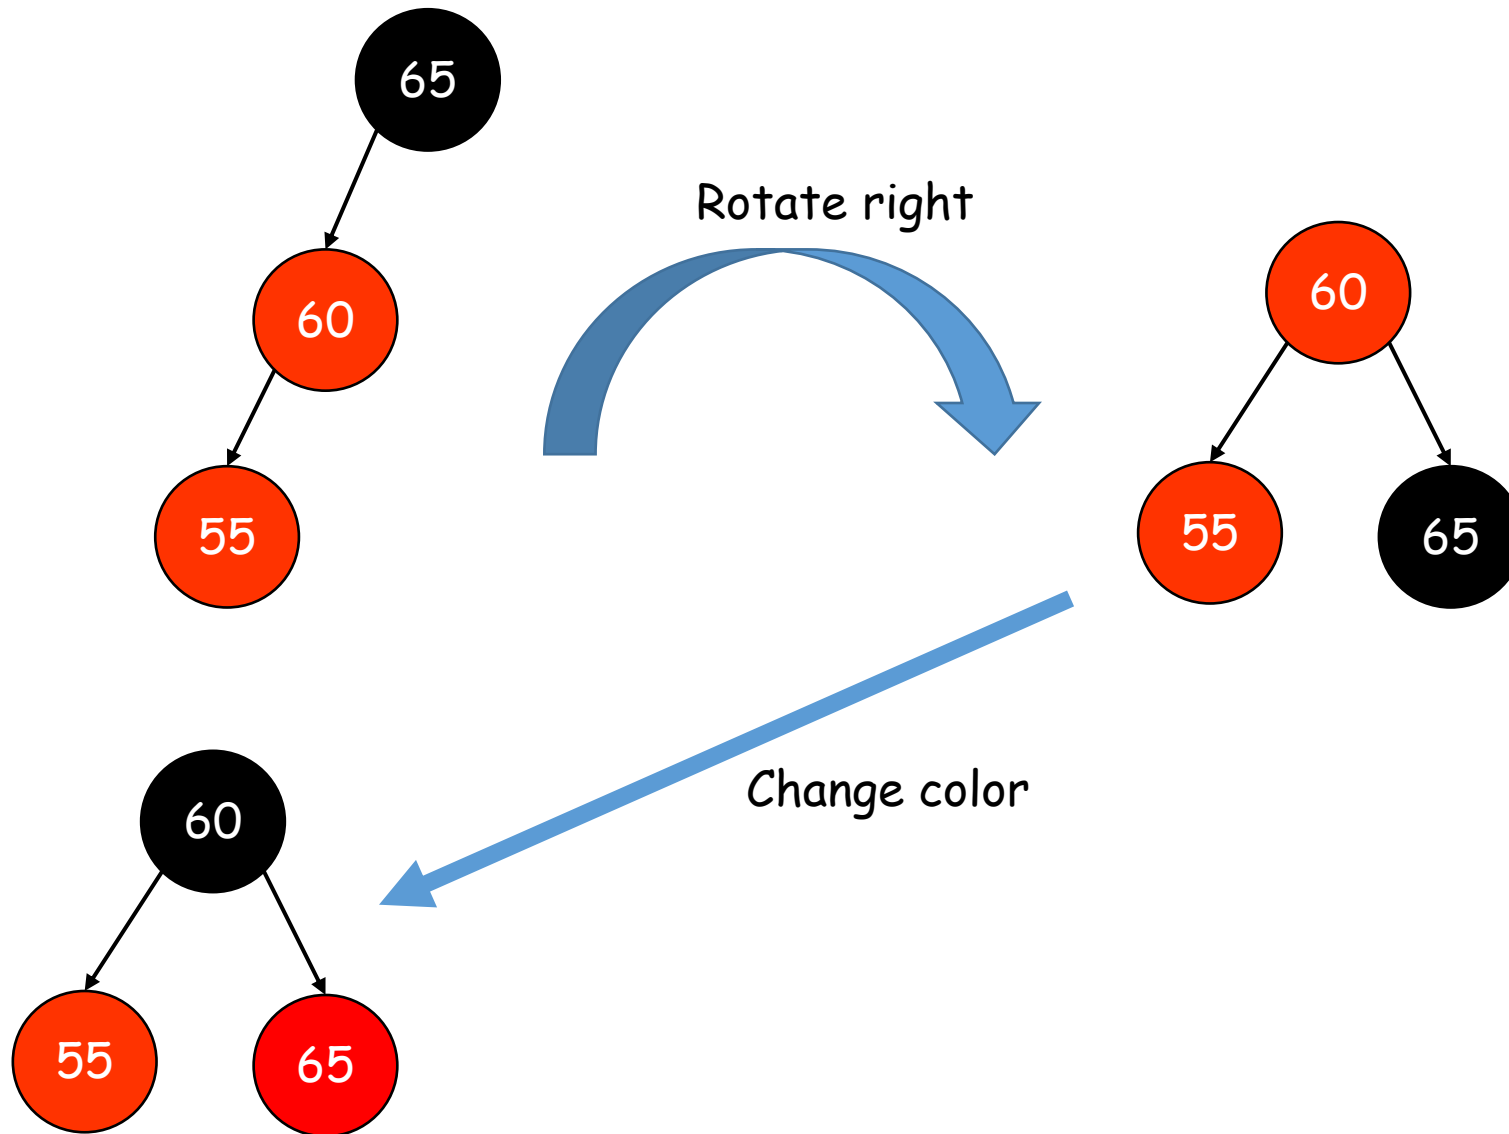
Insert(55)




Insert(55)




Insert(55)



Insert(53) ?



Insert(53)



Problem of the day

Given an array of integers, and x , find how many pairs of elements of the array sum to x , i.e., how many indexes $i \neq j$ are there with $arr[i] + arr[j] = x$? Array is not sorted, and may contain duplicate elements. Solve the problem using balanced BST with RT = $O(n \log n)$. For example, If $arr = \{3, 3, 4, 5, 3, 5, 4\}$ then $howMany(A, 8)$ returns 7.